GlusterFS is an easy-to-use distributed storage system that is very accommodating to your data needs.

GlusterFS isn't just a filesystem. It's a Lego-like tool kit for creating scalable storage systems. With our architecture and hackability, you can create any kind of system you want.

Unified File and Object (UFO)

This feature is unique to GlusterFS – nothing else gives you the same data over this many protocols and access points, proprietary or otherwise.

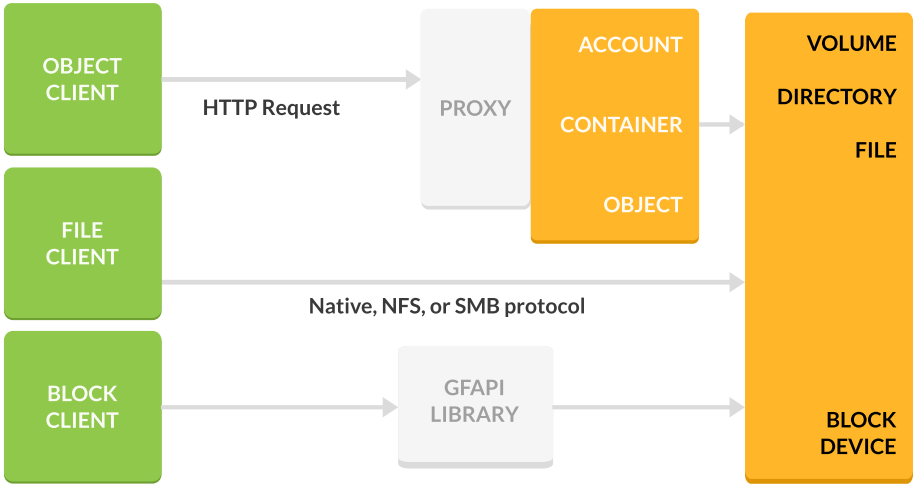

GlusterFS in a nutshell

You're not restricted to keeping your storage system in the same data center. With asynchronous replication you can manage master-slave replicated volumes wherever there is a network connection.

Management

In about five commands, you can have GlusterFS running on multiple nodes, serving up distributed and replicated data. And we run on any disk filesystem that supports extended attributes, so removing GlusterFS doesn't remove your file or data access.

Our CLI allows you to change volume configurations on the fly, without having to take down GlusterFS services. What's more, we've recently collaborated with the oVirt project (ovirt.org) to give users the ability to manage their GlusterFS volumes via a GUI and RESTful integration platform.

Hackability

No one else gives you the same level of hackability as GlusterFS. HekaFS (hekafs.org) and the recently contributed QEMU integration from IBM (introduced in GlusterFS 3.4) all make extensive use of GlusterFS translator APIs. GlusterFS also speaks OpenStack Swift and is POSIX compliant.

More information

Website:
gluster.org

Twitter:
@Glusterorg